

## SYSCALL IMPLEMENTATION

Two syscalls are implemented namely `kernel_device_reader` and `kernel_device_writer`. The reader syscall dequeues data block from the kernel queue, and the writer syscall enqueues data to the kernel queue which is a thread-safe character device.

```
SYSCALL_DEFINE2(kernel_device_reader, char __user *, read_buffer, int, len_buffer)
```

Reader syscall takes in a user-space `char *` and the size of the string which is fixed to 8bytes, the data block dequeued from the kernel queue, is stored in the `char *`, so consumer program 'c' could print the data block in userspace.

Reader syscall is called by the 'c' file.

```
SYSCALL_DEFINE2(kernel_device_writer, char __user *, write_buffer, int, len_buffer)
```

Writer syscall takes in a user-space `char *` and the size of the string which is fixed to 8bytes, the data from this `char*` is enqueued to the kernel queue(character device). so producer program 'p' reads random 8-bytes of the device `/dev/urandom` and then passes this to the writer syscall, which enqueues this on the kernel queue.

Writer syscall is called by the 'p' executable file.

To open the kernel queue(as it is a character device ). `filp_open` is used

```
struct file *filp_open(const char *, int, umode_t);
```

To close the character device `filp_close` is used.

```
int filp_close(struct file *, fl_owner_t id);
```

To read(dequeue) from the kernel queue, `vfs_read` is used.

```
ssize_t vfs_read(struct file *, char __user *, size_t, loff_t *);
```

To write(enqueue) to the kernel queue, `vfs_write` is used.

```
ssize_t vfs_write(struct file *, const char __user *, size_t, loff_t *)
```

All the errors generated during file opening, closing, reading and writing are checked and syscall returns `-EFAULT` and `errno` is set accordingly and the error is printed to stdout using `perror`.

If the syscall is successful 0 is returned.

### Kernel log:

Also when the system call is used, data relevant to the syscall is also printed in the kernel logs. In case of any error, it is printed in the kernel log and the log can be accessed by `dmesg | tail`.

## Kernel Queue

### About the kernel queue:

A kernel level queue is used as a character device, to exchange strings between user-level processes. The device maintains a FIFO queue that can contain a configurable number of strings. Several concurrent user-level processes can read and write to the character device. So the kernel queue is made thread-safe so various user-level processes can read and write to the character device without any race conditions. This is achieved by the use of kernel synchronization primitives like kernel reader-writer semaphores. Some of them are:

```
struct semaphore {
    raw_spinlock_t    lock;
    unsigned int      count;
    struct list_head wait_list;
};

void sema_init(struct semaphore *sem, int val)
void up(struct semaphore *sem);
int down_interruptible(struct semaphore *sem)
```

`Sema_init` is used to initialize the semaphore,

`Up` is used to increment the value of the semaphore (count variable) is similar to `sem_post()`

`Down_interruptible` is used to decrement the value of the semaphore and is similar to `sem_wait()`.

To install the character device(kernel queue):

- 1) Do `"cd new_device"`
- 2) Do `"make"`
- 3) Do `"sudo insmod new_device.ko buffer_size=<size>"` to install the module.

Now our kernel queue is installed by the name of “new\_device” and can be looked up by the command “**lsmod**” which lists all the loaded modules and registration about the device is also printed in the kernel log.

The kernel queue (character device) is installed in /dev/new\_device.

If the kernel fails to install and load, details are printed in the kernel log and can be checked by Dmesg | tail.

```
[kern@admin9 new_device]$ lsmod
Module                Size  Used by
new_device            16384  0
vmwgfx                380928  1
intel_rapl_msr         20480  0
intel_rapl_common      28672  1 intel_rapl_msr
vmw_balloon            24576  0
crct10dif_pclmul       16384  1
crc32_pclmul           16384  0
```

The kernel implements the following functions,

```
.open = &open_device,
.read = &dequeue_data,
.write = &enqueue_data,
.release = &close_device
```

To remove the loaded module, do “**rmmod new\_device**”.

### Modified files:

- 1)arch/x86/entry/syscalls/syscall\_64.tbl
- 2)kernel/sys.c
- 3) /dev/new\_device

### Diff files/ Patch file:

Diff.txt and Patch\_file.patch is attached along with the code.

Commands used:

Git diff > diff.txt

Git diff > patch\_file.patch

## **Producer.c and consumer.c**

Makefile to compile both producer.c and consumer.c is provided. Do “make” which makes two executable files namely p and c.

To run these files do :

**Sudo ./c first  
then sudo ./p**

The implemented syscalls and the kernel queue loaded can be tested using these files, which implements the producer-consumer problem. All the errors are checked and errors are printed out by perror().

Producer.c reads random 8-bytes of the device /dev/urandom and passes them to the kernel via a system call, kernel\_device\_write. This system call enqueues the data sent by producer.c to the kernel queue.

And consumer.c calls the kernel\_device\_read system call which dequeues the data from the kernel and this block of data(string) is sent back to the consumer where this string is printed.